

Operating systems and concurrency (B10)

David Kendall

Northumbria University

- Semaphores provide an **unstructured** synchronisation primitive
- Can lead to problems:
 - Accidental release
 - Deadlock
 - Starvation
 - ...
- Can be difficult to detect and debug
- **Monitors** and **condition variables** offer one approach to a more structured synchronisation mechanism
- Not supported directly in uC/OS-II but widely used
- E.g. POSIX threads (Pthreads)


Monitor

- A high-level abstraction that provides a convenient and effective mechanism for process synchronization
- Only one process may be active within the monitor at a time – **mutual exclusion**

```
monitor monitor-name {  
    // shared variable declarations  
  
    procedure P1 (...) { ... }  
  
    ...  
  
    procedure Pn (...) { ... }  
  
    Initialization code (...) { ... }  
}
```

- Proposed independently by Per Brinch Hansen and Tony Hoare in 1973/74

Schematic view of a monitor



- Thread executing a monitor operation must hold the monitor mutex
- Other threads wanting to execute a monitor operation are queued on the mutex until the executing thread finishes and releases the mutex


(from [SGG10, p.247])

Condition variables

- So a monitor provides mutual exclusion
- **Condition variables** introduced to allow **signalling**
- Three operations allowed on condition variable, `cv`:
 - **`wait(cv)`** – block until another thread calls **signal** or **broadcast** on `cv`
 - **`signal(cv)`** – wake up *one* thread waiting on `cv`
 - **`broadcast(cv)`** – wake up *all* threads waiting on `cv`
- In Pthreads the CV type is `pthread_cond_t`
 - Use `pthread_cond_init()` to initialise
 - `pthread_cond_wait(&cv, &mutex);`
 - `pthread_cond_signal(&cv);`
 - `pthread_cond_broadcast(&cv);`

(adapted from slides by Matt Welsh, Harvard University, 2009)

Monitor with condition variables



- A thread that waits on a condition **releases the monitor mutex**
- When it is signalled, it must **reacquire the mutex** before continuing


(from [SGG10, p.248])

Hoare vs Mesa Monitor Semantics

- The monitor `signal()` operation can have two different meanings:
- Hoare monitors (1974)
 - `signal(cv)` means to run the waiting thread immediately
 - Effectively “hands the lock” to the thread just signaled.
 - Causes the signalling thread to block
- Mesa monitors (Xerox PARC, 1980)
 - `signal(cv)` puts waiting thread back onto the “ready queue” for the monitor
 - But, signaling thread keeps running.
 - Signaled thread doesn’t get to run until it can acquire the lock.
 - This is what we almost always use, eg Pthreads, Java, C#, etc
- What’s the practical difference?
 - In Hoare-style semantics, the “condition” that triggered the `signal()` will always be true when the awoken thread runs
 - For example, that the buffer is now no longer empty
 - In Mesa-style semantics, awoken thread has to recheck the condition
 - Since another thread might have beaten it to the punch

(adapted from slides by Matt Welsh, Harvard University, 2009)

Dining philosophers



- Another classic synchronisation problem
- Introduced here to illustrate how to build a monitor using **Pthreads**
- Philosophers **think**, get **hungry**, then **eat**, ... repeatedly, ... that's all
- To eat, a philosopher must have **two** chopsticks
- Spot the deadlock possibility ... what about starvation? ... oh how we laughed!
- See (dining philosophers monitor in Pthreads)


```
#include <assert.h>
#include <sys/types.h>
#include <pthread.h>
#include <stdio.h>
#include <stdlib.h>
#include <string.h>
#include "dpmonitor.h"

/* ***** Type declarations ***** */

typedef enum {THINKING, HUNGRY, EATING} state_t;

/* ***** Local function prototypes ***** */

static void eatIfOk(int i);
static int leftNghbr(int i);
static int rightNghbr(int i);
```

```
/* ***** Monitor variables ***** */
```

```
static pthread_mutex_t dpMutex;
```

```
static pthread_cond_t okToEat[N_PHIL];
```

```
static state_t state[N_PHIL];
```

```
/* ***** Monitor function definitions ***** */
```

```
void dpInit(void) {  
    int rc;  
  
    // Initialise the monitor mutex  
    rc = pthread_mutex_init(&dpMutex, NULL);  
    assert(rc == 0);  
  
    // Initialise the state and the condition variables  
    for (int i=0; i<N_PHIL; i+=1) {  
        state[i] = THINKING;  
        rc = pthread_cond_init(&okToEat[i], NULL);  
        assert(rc == 0);  
    }  
}
```

```
void dpPickup(int i) {  
    pthread_mutex_lock(&dpMutex);  
    state[i] = HUNGRY;  
    eatIfOk(i);  
    while (state[i] != EATING) {  
        pthread_cond_wait(&okToEat[i], &dpMutex);  
    }  
    pthread_mutex_unlock(&dpMutex);  
}
```

```
void dpPutdown(int i) {  
    pthread_mutex_lock(&dpMutex);  
    state[i] = THINKING;  
    eatIfOk(rightNghbr(i));  
    eatIfOk(leftNghbr(i));  
    pthread_mutex_unlock(&dpMutex);  
}
```

```
/* ***** Local function definitions ***** */

static void eatIfOk (int i) {
    if ((state[i] == HUNGRY) &&
        (state[rightNghbr(i)] != EATING) &&
        (state[leftNghbr(i)] != EATING)) {
        state[i] = EATING;
        pthread_cond_signal(&okToEat[i]);
    }
}

static int leftNghbr(int i) {
    return ((i+(N_PHIL-1)) % N_PHIL);
}

static int rightNghbr(int i) {
    return ((i+1) % N_PHIL);
}
```

Acknowledgements

- **[SGG10]** Silberschatz, A., Galvin, P., Gagne, G., *Operating System Concepts* (8th edition), Wiley, 2010
- **Lawrence Livermore Pthreads tutorial**
- Welsh, M., *Semaphores, Condition Variables and Monitors*, Lecture slides, Harvard University, 2009 (**local copy**)